

Relational Algebra

CS430/630
Lecture 2

Slides based on "Database Management Systems" 3rd ed, Ramakrishnan and Gehrke

Relational Query Languages

- ▶ Query languages:
 - ▶ Allow manipulation and retrieval of data from a database
- ▶ Relational model supports simple, powerful QLs:
 - ▶ Strong formal foundation based on logic
 - ▶ Allows for much optimization
- ▶ Query Languages != programming languages
 - ▶ QLs not intended to be used for complex calculations
 - ▶ QLs support easy, efficient access to large data sets

Formal Relational Query Languages

- ▶ Two languages form the basis for SQL:
 - ▶ **Relational Algebra:**
 - ▶ operational
 - ▶ useful for representing execution plans
 - ▶ very relevant as it is used by query optimizers!
 - ▶ **Relational Calculus:**
 - ▶ Lets users describe the result, NOT how to compute it - declarative
 - ▶ We will focus on relational algebra

Preliminaries

- ▶ A query is applied to **relation instances**, and the result of a query is also a relation instance
 - ▶ Schemas of input relations for a query are **fixed**
 - ▶ The schema for the **result** of a given query is determined by operand schemas and operator type
- ▶ Each operation returns a relation
 - ▶ operations can be **composed**!
 - ▶ Well-formed expression: a relation, or the results of a relational algebra operation on one or two relations

Relational Algebra

- ▶ Basic operations:
 - ▶ **Selection** σ Selects a subset of rows from relation
 - ▶ **Projection** π Deletes unwanted columns from relation
 - ▶ **Cross-product** \times Allows us to combine several relations
 - ▶ **Join** \bowtie Combines several relations using conditions
 - ▶ **Division** \div A bit more complex, will cover later on
 - ▶ **Set-difference** $-$ **Union** \cup **Intersection** \cap
 - ▶ **Renaming** ρ Helper operator, does not derive new result, just renames relations and fields

$$\rho(R(F), E)$$
 - ▶ F contains $oldname \rightarrow newname$ pairs

Example Schema

Sailors				Boats		
<u>sid</u>	sname	rating	age	<u>bid</u>	name	color
22	dustin	7	45.0	101	interlake	red
31	lubber	8	55.5	103	clipper	green
58	rusty	10	35.0			

Reserves		
<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Relation Instances Used

S1				S2			
sid	sname	rating	age	sid	sname	rating	age
22	dustin	7	45.0	28	yuppy	9	35.0
31	lubber	8	55.5	31	lubber	8	55.5
58	rusty	10	35.0	44	guppy	5	35.0
				58	rusty	10	35.0

R1		
sid	bid	day
22	101	10/10/96
58	103	11/12/96

Projection

- Unary operator
 - Deletes (projects out) attributes that are not in *projection list*
- $$\pi_{attr1, attr2, \dots} relation$$
- Result Schema* contains the attributes in the projection list
 - With the same names that they had in the input relation
 - Projection operator has to eliminate *duplicates*!
 - Real systems typically do not do so by default
 - Duplicate elimination is *expensive*! (sorting)
 - User must explicitly ask for duplicate eliminations (**DISTINCT**)

Projection Example

S2			
sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating
28	yuppy	9
31	lubber	8
44	guppy	5
58	rusty	10

$\pi_{sname, rating}(S2)$

Selection

- Unary Operator
 - Selects rows that satisfy *selection condition*
- $$\sigma_{condition} relation$$
- Condition* contains constants and attributes from relation
 - Evaluated for each *individual* tuple
 - May use logical connectors AND (^), OR (v), NOT (-)
 - No duplicates in result! Why?
 - Result Schema* is identical to schema of the input relation

Selection Example

S2			
sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

sid	sname	rating
28	yuppy	9
58	rusty	10

$\sigma_{rating > 8}(S2)$

Selection and Projection $\pi_{sname, rating}(\sigma_{rating > 8}(S2))$

Cross-Product

- Binary Operator
- $$R \times S$$
- Each row of relation R is paired with each row of S
 - Result Schema* has one field per field of R and S
 - Field names 'inherited' when possible

Cross-Product Example

S1	sid	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

R1	sid	bid	day
	22	101	10/10/96
	58	103	11/12/96

C=S1 X R1

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Conflict: Both R and S have a field called **sid**

Cross-Product + Renaming Example

C

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Renaming operator $\rho(C(1 \rightarrow \text{sid1}, 5 \rightarrow \text{sid2}), S1 \times R1)$

Condition Join (Theta join)

$$R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$$

- Result Schema same as that of cross-product

Condition Join Theta-join Example

S1 X R1

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

S1 $\bowtie_{sid1 < R1.sid}$ R1

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

Equi-Join

- A special case of condition join where the condition contains only **equalities**

$$R \bowtie_{R.attr1 = S.attr2} S$$

- Result Schema similar to cross-product, but only one copy of fields for which equality is specified.

Equi-Join Example

S1 X R1

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

S1 $\bowtie_{sid} R1$

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

Natural Join

- Equijoin on **all** common fields

$$R \bowtie S$$

- Common fields are **NOT** duplicated in the result

Union, Intersection, Set-Difference

- All of these operations take two input relations, which must be **union-compatible**
 - Same number of fields.
 - Corresponding fields have the same domain (type)
- What is the *schema* of result?

Union Example

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

Intersection Example

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

Set-Difference Example

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	sname	rating	age
22	dustin	7	45.0

$S1 - S2$